

Socket Programming

- ❑ What is a socket?

- ❑ Using sockets

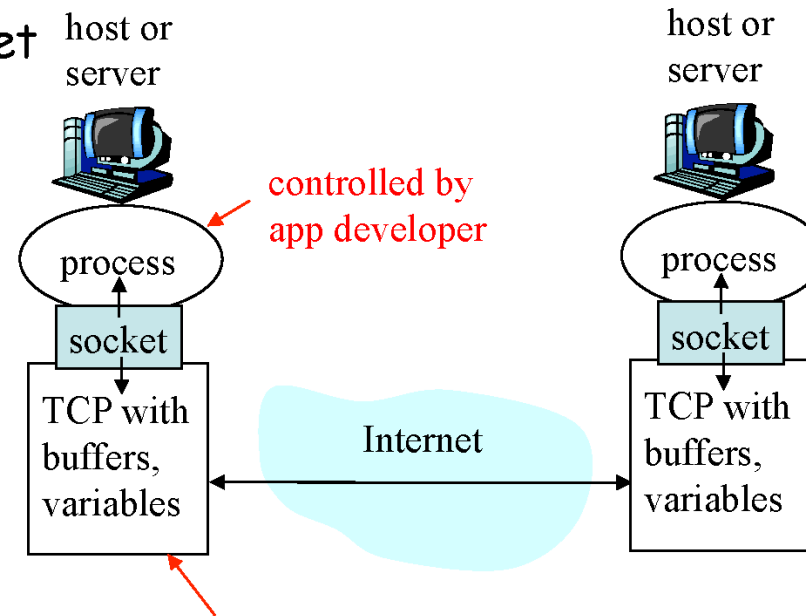
- Types (Protocols)
- Associated functions
- Styles

- Socket programming reference:

- TCP/IP 소켓 프로그래밍 - C버전, Michael J. Donahoo, (박준철 번역), 사이텍미디어

What is a socket?

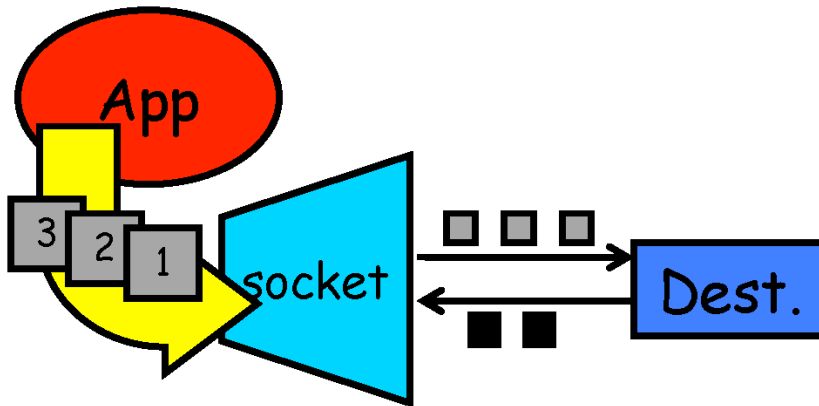
- ❑ An interface between application and network
 - The application creates a socket
 - The socket *type* dictates the style of communication
 - reliable vs. best effort
 - connection-oriented vs. connectionless
- ❑ Once configured, the application can
 - pass data to the socket for network transmission
 - receive data from the socket (transmitted through the network by some other host)



Two essential types of sockets

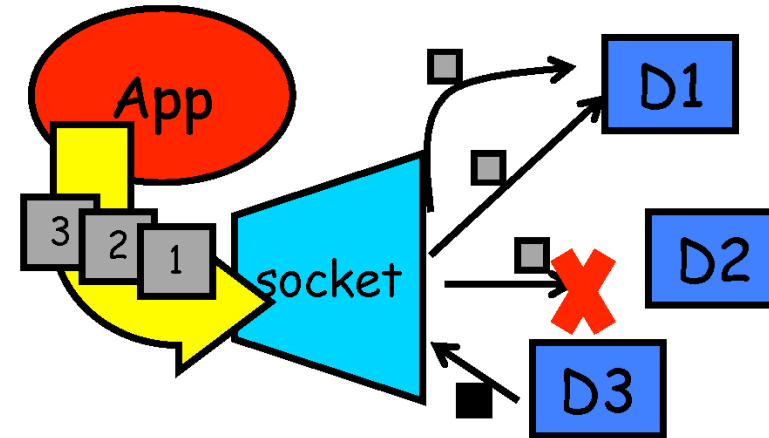
□ SOCK_STREAM

- a.k.a. **TCP**
- reliable delivery
- in-order guaranteed
- connection-oriented
- bidirectional



□ SOCK_DGRAM

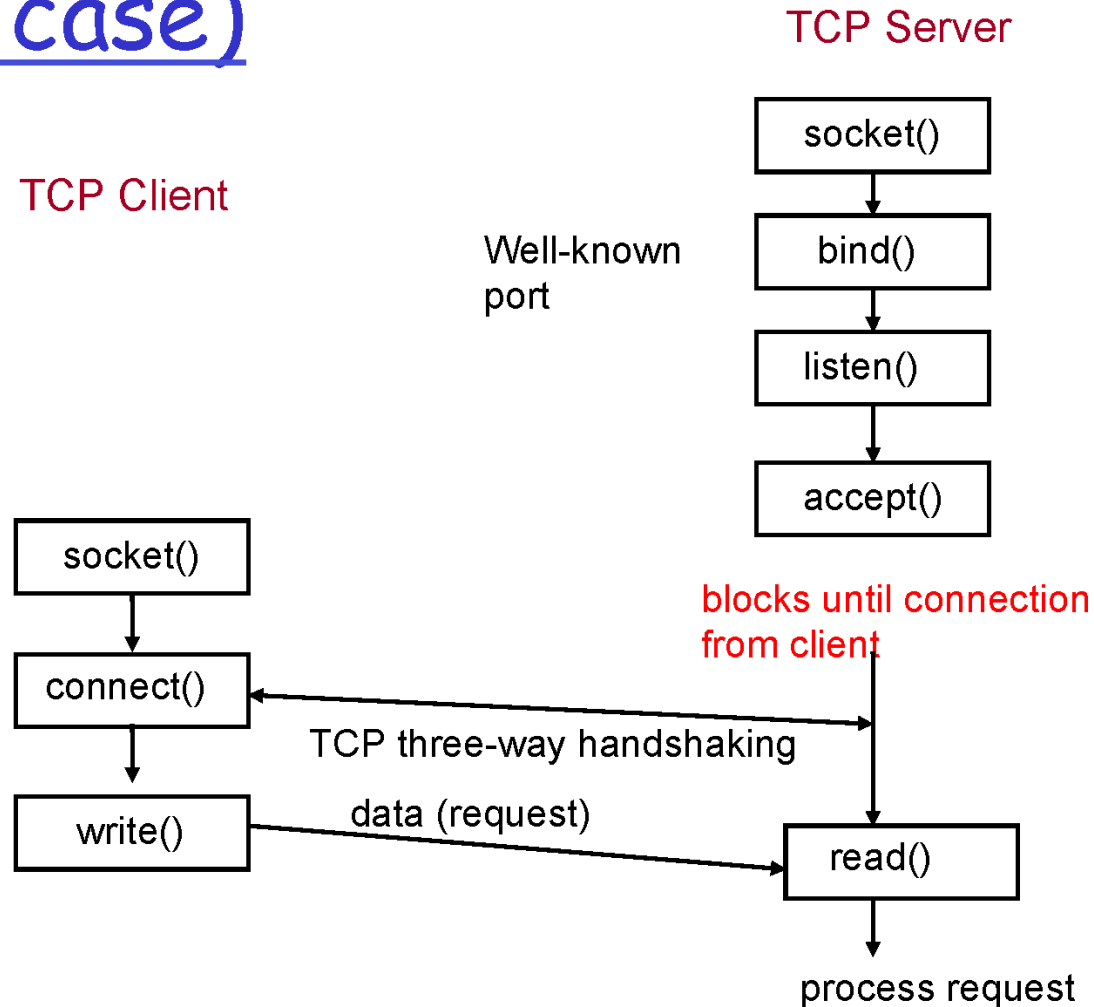
- a.k.a. **UDP**
- unreliable delivery
- no order guarantees
- no notion of "connection" - app indicates dest. for each packet
- can send or receive



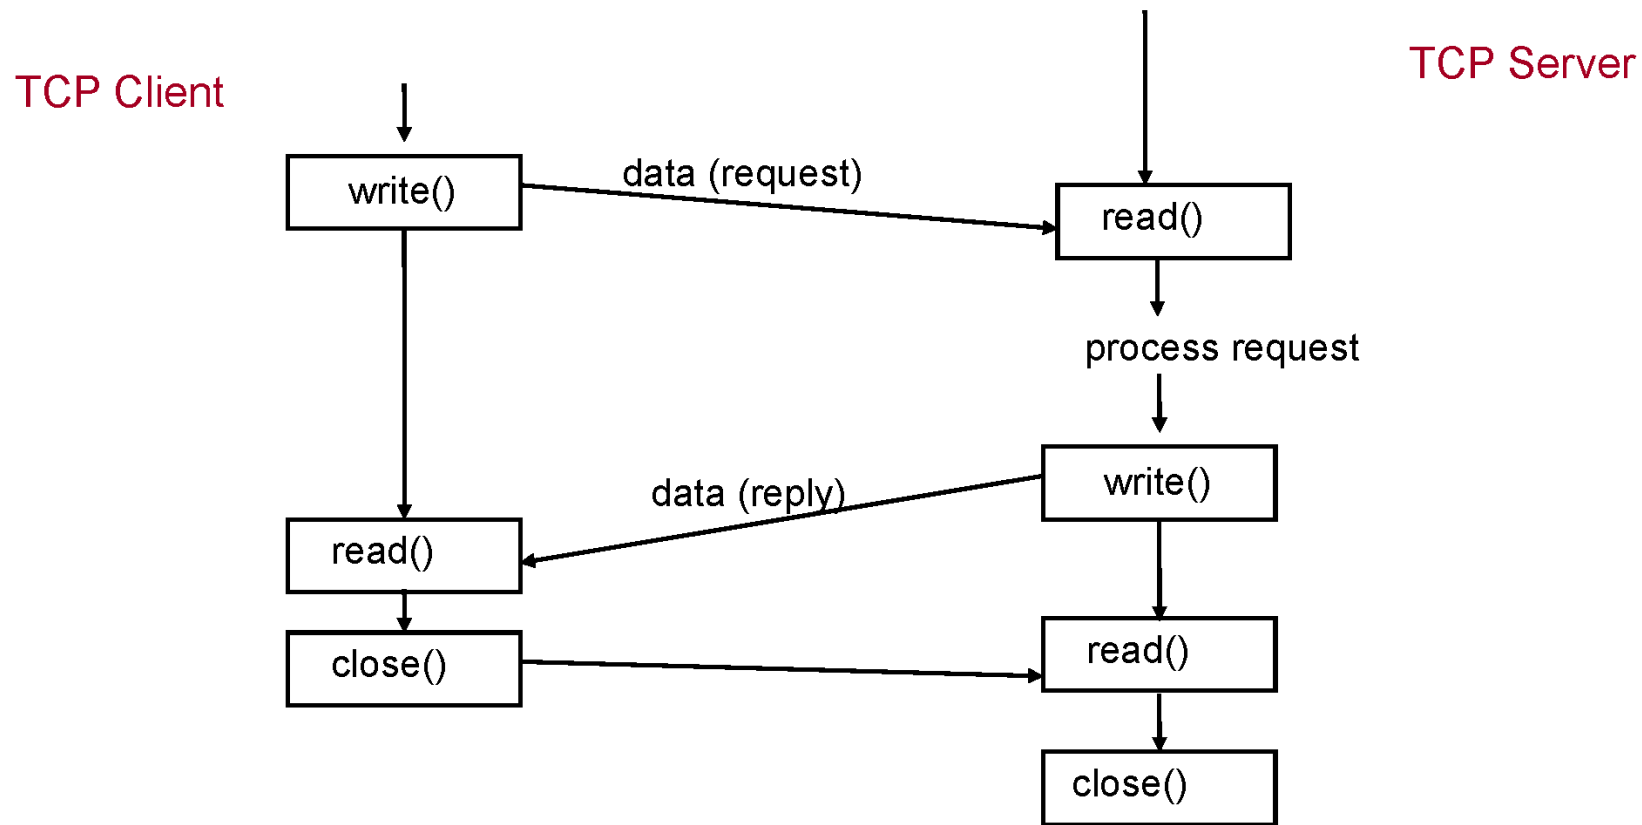
Sockets API

- ❑ Creation and Setup
- ❑ Establishing a Connection (TCP)
- ❑ Sending and Receiving Data
- ❑ Tearing Down a Connection (TCP)

Big picture: Socket Functions (TCP case)



Big picture: Socket Functions (TCP case) cont.



Socket Creation and Setup

- ❑ Include file `<sys/socket.h>`
- ❑ **Create** a socket
 - `int socket (int domain, int type, int protocol);`
 - Returns file descriptor or -1.
- ❑ **Bind** a socket to a local IP address and port number
 - `int bind (int sockfd, struct sockaddr* myaddr, int addrlen);`
- ❑ Put socket into passive state (**wait for connections** rather than initiate a connection).
 - `int listen (int sockfd, int backlog);`
- ❑ **Accept** connections
 - `int accept (int sockfd, struct sockaddr* cliaddr, int* addrlen);`
 - Returns file descriptor or -1.

Function: socket

```
int socket (int domain, int type, int
            protocol) ;
```

❑ Create a socket.

- Returns file descriptor or -1. Also sets `errno` on failure.
- domain: protocol family (same as address family)
 - `PF_INET` for IPv4 (typicall used)
 - other possibilities: `PF_INET6` (IPv6), `PF_UNIX` or `PF_LOCAL` (Unix socket), `PF_ROUTE` (routing)
- type: style of communication
 - `SOCK_STREAM` for TCP (with `PF_INET`)
 - `SOCK_DGRAM` for UDP (with `PF_INET`)
- protocol: protocol within family
 - Typically set to 0
 - `getprotobyname()`, `/etc/protocols` for list of protocols

Function: bind

```
int bind (int sockfd, struct sockaddr*  
         myaddr, int addrlen) ;
```

- ❑ Bind a socket to a local IP address and port number.
 - Returns 0 on success, -1 and sets `errno` on failure.
 - `sockfd`: socket file descriptor (returned from `socket`)
 - `myaddr`: includes IP address and port number
 - IP address: set by kernel if value passed is `INADDR_ANY`, else set by caller
 - port number: set by kernel if value passed is 0, else set by caller
 - `addrlen`: length of address structure
 - = `sizeof (struct sockaddr_in)`

Function: listen

```
int listen (int sockfd, int backlog) ;
```

- ❑ Put socket into passive state (wait for connections rather than initiate a connection).
 - Returns 0 on success, -1 and sets errno on failure.
 - sockfd: socket file descriptor (returned from socket)
 - backlog: bound on length of unaccepted connection queue (connection backlog); kernel will cap, thus better to set high
- Listen is non-blocking: returns immediately

Function: accept

```
int accept (int sockfd, struct sockaddr*  
            cliaddr, int* addrlen);
```

❑ Accept a new connection.

- Returns file descriptor or -1. Also sets `errno` on failure.
- `sockfd`: socket file descriptor (returned from `socket`)
- `cliaddr`: IP address and port number of client (returned from call)
- `addrlen`: length of address structure = pointer to `int` set to `sizeof (struct sockaddr_in)`

❑ Accept is blocking

- Waits for connection before returning

Sockets API

- ❑ Creation and Setup
- ❑ Establishing a Connection (TCP)
- ❑ Sending and Receiving Data
- ❑ Tearing Down a Connection (TCP)

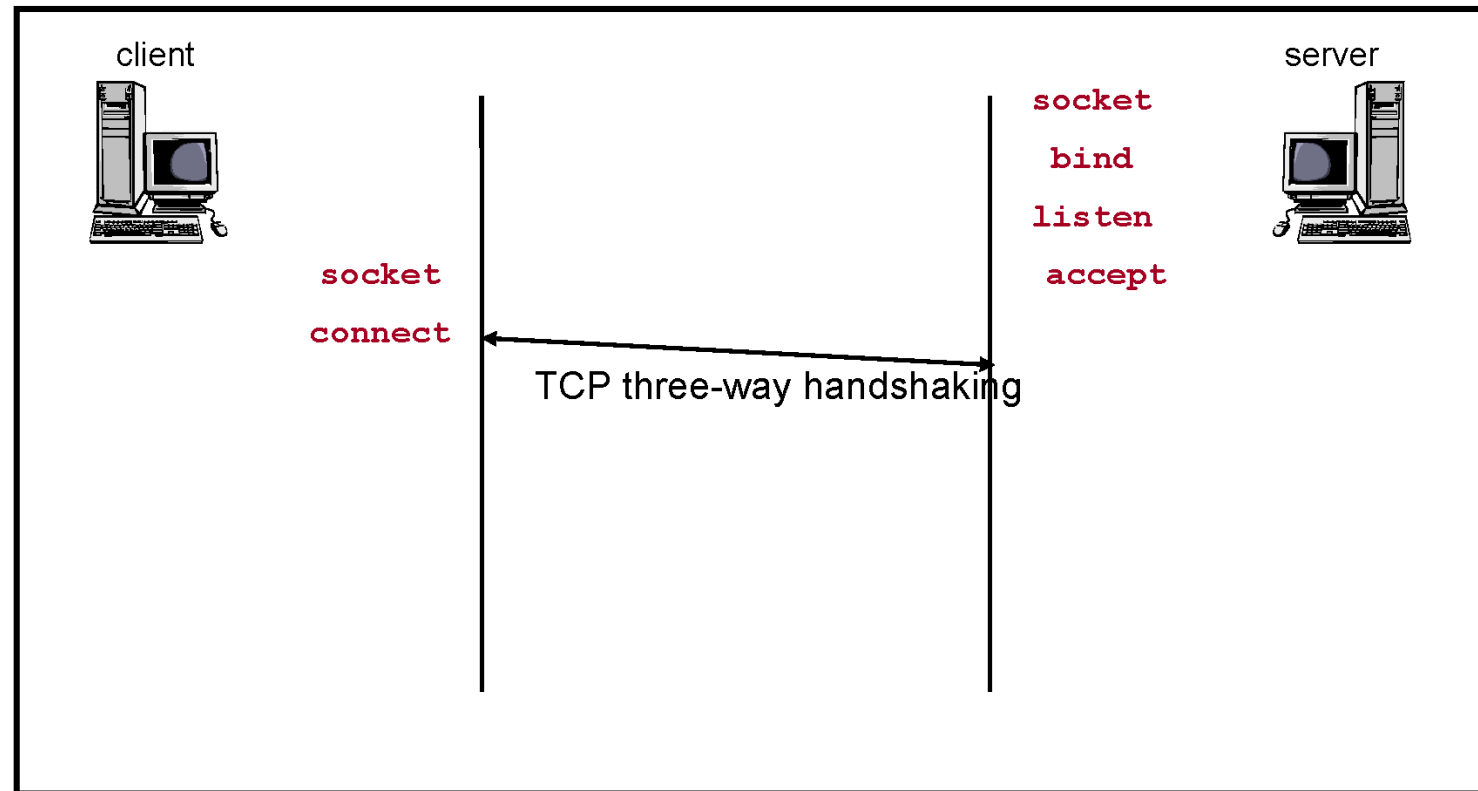
Function: connect

```
int connect (int sockfd, struct sockaddr*  
             servaddr, int addrlen);
```

❑ Connect to another socket.

- Returns 0 on success, -1 and sets errno on failure.
 - sockfd: socket file descriptor (returned from socket)
 - servaddr: IP address and port number of **server**
 - addrlen: length of address structure
 - = sizeof (struct sockaddr_in)
- Connect is **blocking**

Recap: TCP socket connection setup



Sample code: server

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <string.h>
#include <sys/types.h>
#include <netinet/in.h>
#include <sys/socket.h>
#include <sys/wait.h>
#define PORT 3490
#define BACKLOG 10      /* how many pending
                        connections queue
                        will hold */
```

server

```
main()
{
    int sockfd, new_fd;      /* listen on sock_fd, new
                             connection on new_fd */
    struct sockaddr_in my_addr; /* my address */
    struct sockaddr_in their_addr; /* connector addr */
    int sin_size;

    if ((sockfd = socket(PF_INET, SOCK_STREAM, 0)) == -1) {
        perror("socket");
        exit(1);
    }
```


server

```
my_addr.sin_family = AF_INET;
my_addr.sin_port = htons(MYPORT); /* short, network
                                   byte order */
my_addr.sin_addr.s_addr = htonl(INADDR_ANY);
/* INADDR_ANY allows clients to connect to any one of
the host's IP address */

if (bind(sockfd, (struct sockaddr *)&my_addr,
        sizeof(struct sockaddr)) == -1) {
    perror("bind");
    exit(1);
}
```

- The Internet-specific:
 - struct sockaddr_in {
 - short sin_family;
 - u_short sin_port;
 - struct in_addr sin_addr;
 - };
 - sin_family = AF_INET
 - sin_port: port # (0-65535)
 - sin_addr: IP-address

server

```
if (listen(sockfd, BACKLOG) == -1) {
    perror("listen");
    exit(1);
}
while(1) { /* main accept() loop */
    sin_size = sizeof(struct sockaddr_in);
    if ((new_fd = accept(sockfd, (struct sockaddr*)
                        &their_addr, &sin_size)) == -1) {
        perror("accept");
        continue;
    }
    printf("server: got connection from %s\n",
          inet_ntoa(their_addr.sin_addr));
}
```

client

```
if ((sockfd = socket (PF_INET, SOCK_STREAM, 0)) == -1) {  
    perror ("socket");  
    exit (1);  
}  
  
their_addr.sin_family = AF_INET;  
their_addr.sin_port = htons (Server_Portnumber);  
their_addr.sin_addr = htonl (Server_IP_address);  
  
if (connect (sockfd, (struct sockaddr*)&their_addr,  
            sizeof (struct sockaddr)) == -1) {  
    perror ("connect");  
    exit (1);  
}
```

Sockets API

- ❑ Creation and Setup
- ❑ Establishing a Connection (TCP)
- ❑ **Sending and Receiving Data**
- ❑ Tearing Down a Connection (TCP)

Functions: write

```
int write (int sockfd, char* buf, size_t
          nbytes) ;
```

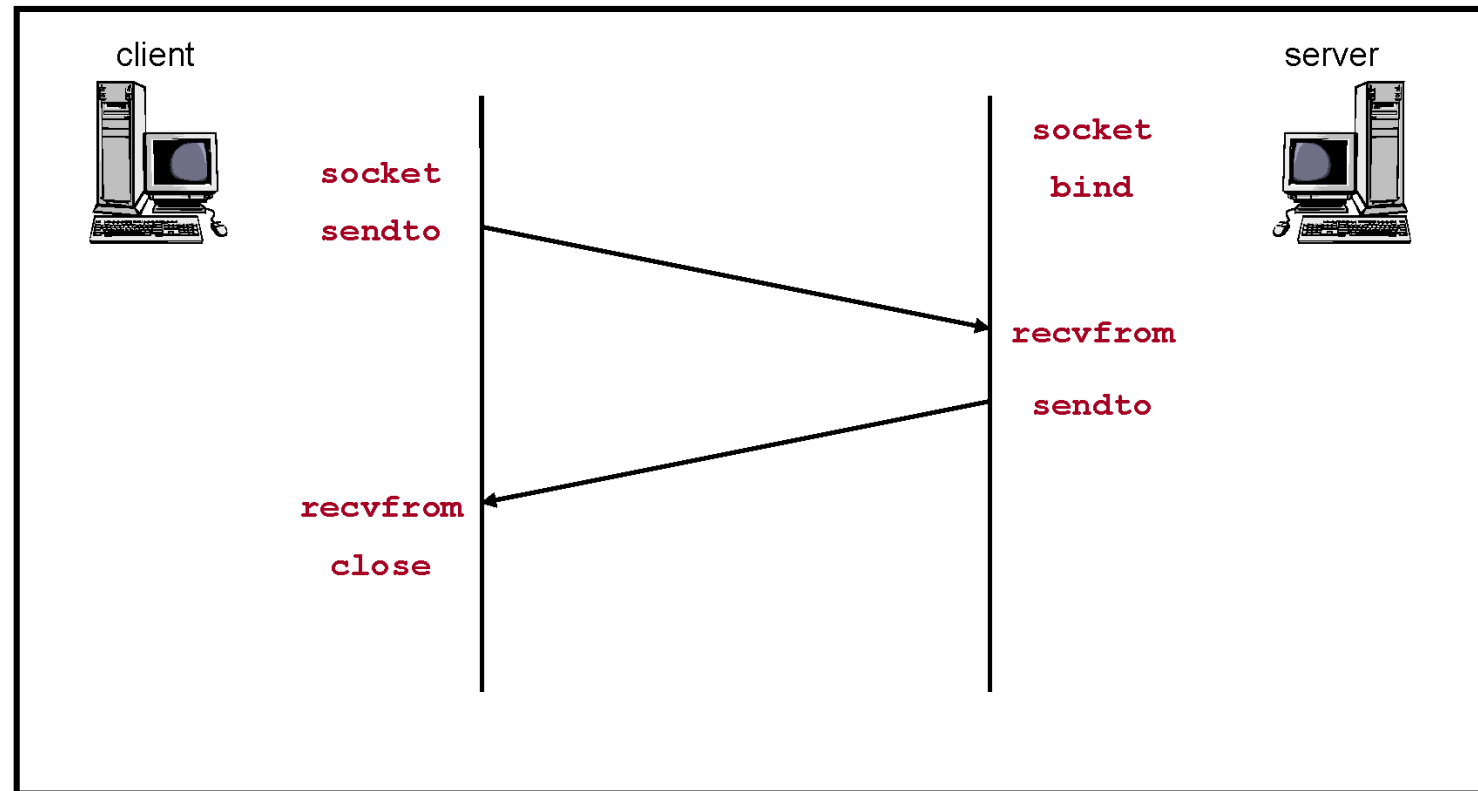
- ❑ Write data to a stream (TCP).
 - Returns number of bytes written or -1. Also sets `errno` on failure.
 - `sockfd`: socket file descriptor (returned from `socket`)
 - `buf`: data buffer
 - `nbytes`: number of bytes to try to write
- ❑ `write` is blocking; returns only after data is sent

Functions: read

```
int read (int sockfd, char* buf, size_t
         nbytes) ;
```

- ❑ Read data from a stream (TCP).
 - Returns number of bytes read or -1. Also sets `errno` on failure.
 - Returns 0 if socket closed.
 - `sockfd`: socket file descriptor (returned from `socket`)
 - `buf`: data buffer
 - `nbytes`: number of bytes to try to read
 - read is **blocking**; returns only after data is received

Big picture: UDP Socket Functions



Sockets API

- ❑ Creation and Setup
- ❑ Establishing a Connection (TCP)
- ❑ Sending and Receiving Data
- ❑ Tearing Down a Connection (TCP)

Function: close

```
int close (int sockfd) ;
```

- ❑ When finished using a socket, the socket should be closed:
 - returns 0 if successful, -1 if error
 - sockfd: the file descriptor (socket being closed)
- ❑ Closing a socket
 - frees up the port used by the socket
 - closes a connection (for SOCK_STREAM)

Tip: Release of ports

- ❑ Sometimes, a "rough" exit from a program (e.g., ctrl-c) does not properly free up a port
- ❑ Eventually (after a few minutes), the port will be freed
- ❑ To reduce the likelihood of this problem, include the following code:
 - `#include <signal.h>`
 - `void cleanExit(){exit(0);}`
 - in socket code:
 - `signal(SIGTERM, cleanExit);`
 - `signal(SIGINT, cleanExit);`

Chapter 3 outline

- ❑ 3.1 Transport-layer services
- ❑ 3.2 Multiplexing and demultiplexing
- ❑ 3.3 Connectionless transport: UDP
- ❑ 3.4 Principles of reliable data transfer
- ❑ 3.5 Connection-oriented transport: TCP
 - segment structure
 - reliable data transfer
 - flow control
 - connection management
- ❑ 3.6 Principles of congestion control
- ❑ 3.7 TCP congestion control

Multiplexing/demultiplexing

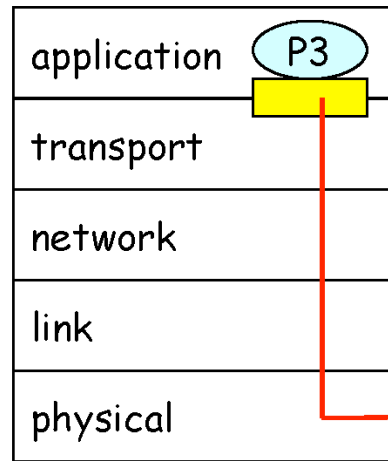
Demultiplexing at rcv host:

delivering received segments
to correct socket

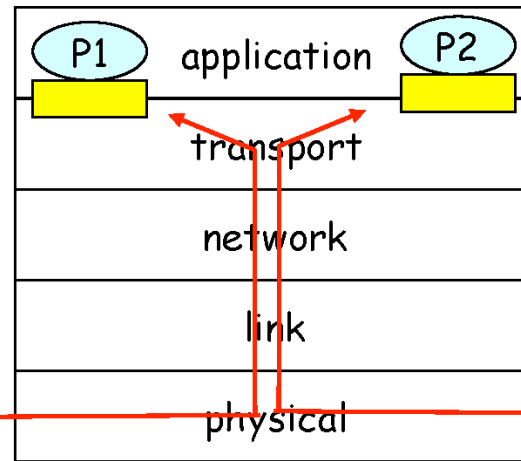
Multiplexing at send host:

gathering data from multiple
sockets, enveloping data with
header (later used for
demultiplexing)

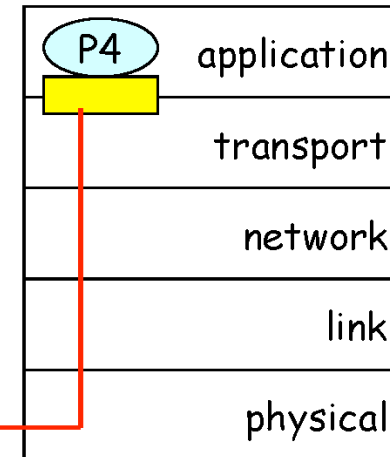
 = socket  = process



host 1



host 2



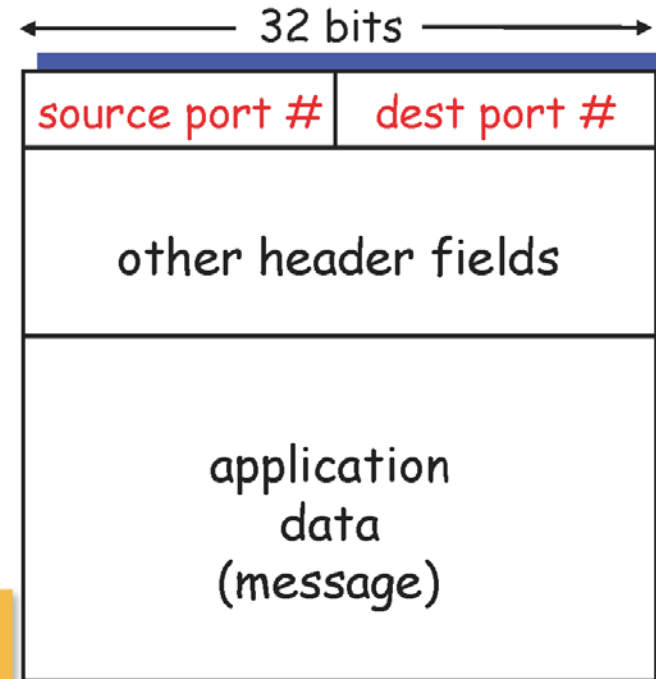
host 3

How demultiplexing works

- ❑ **host receives IP datagrams**
 - each datagram has source IP address, destination IP address
 - each datagram carries 1 transport-layer segment
 - each segment has source, destination port number
- ❑ **host uses IP addresses & port numbers to direct segment to appropriate socket**

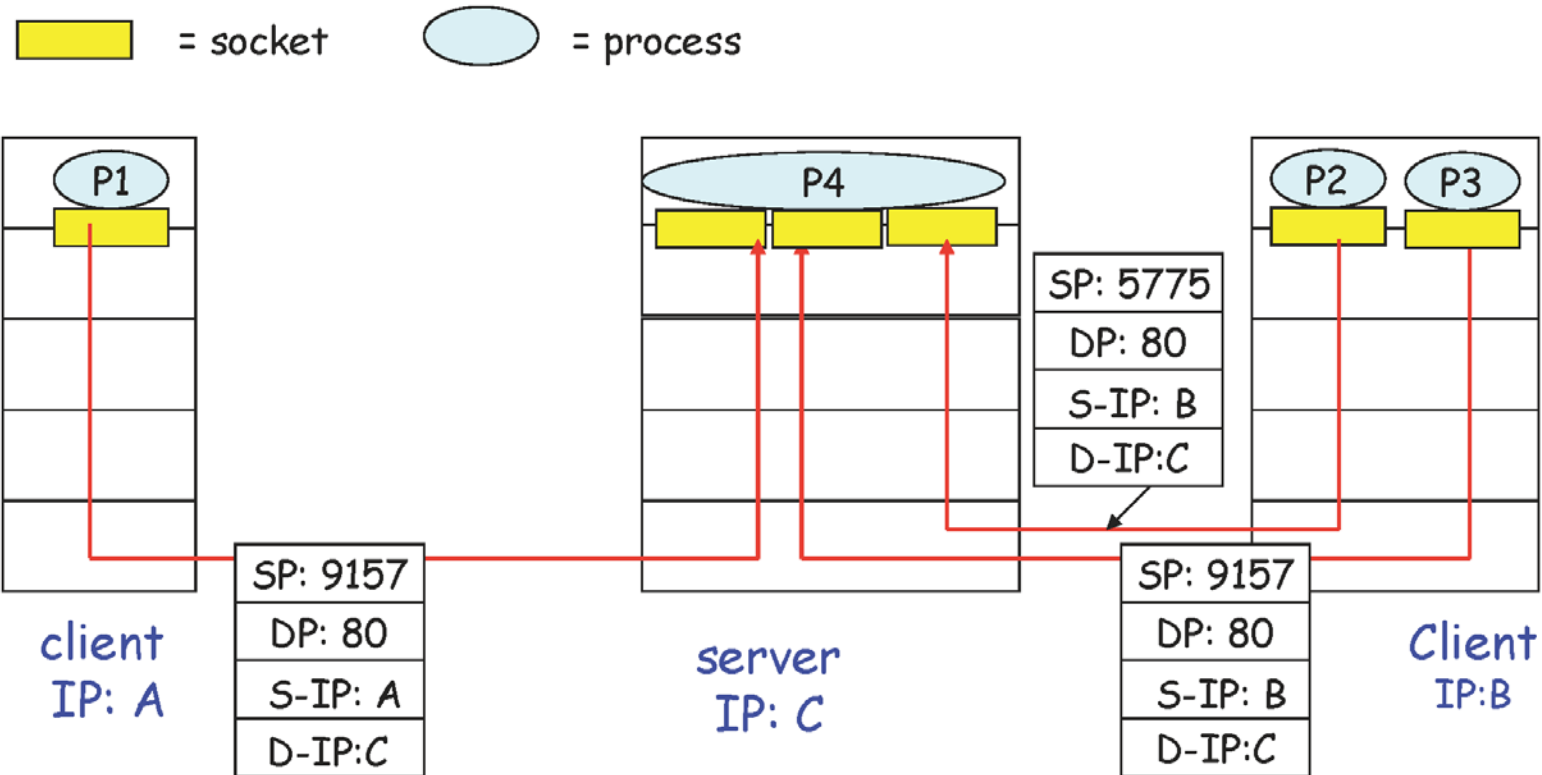
Analogous to airport shuttles

Shuttles MUX passengers and take them to downtown -- DeMUX at different locations



TCP/UDP segment format

Connection-oriented demux: Threaded Web Server



Modify the airport shuttle analogy
to distinguish between UDP and TCP

Chapter 3 outline

- ❑ 3.1 Transport-layer services
- ❑ 3.2 Multiplexing and demultiplexing
- ❑ 3.3 Connectionless transport: UDP
- ❑ 3.4 Principles of reliable data transfer
- ❑ 3.5 Connection-oriented transport: TCP
 - segment structure
 - reliable data transfer
 - flow control
 - connection management
- ❑ 3.6 Principles of congestion control
- ❑ 3.7 TCP congestion control

UDP: User Datagram Protocol [RFC 768]

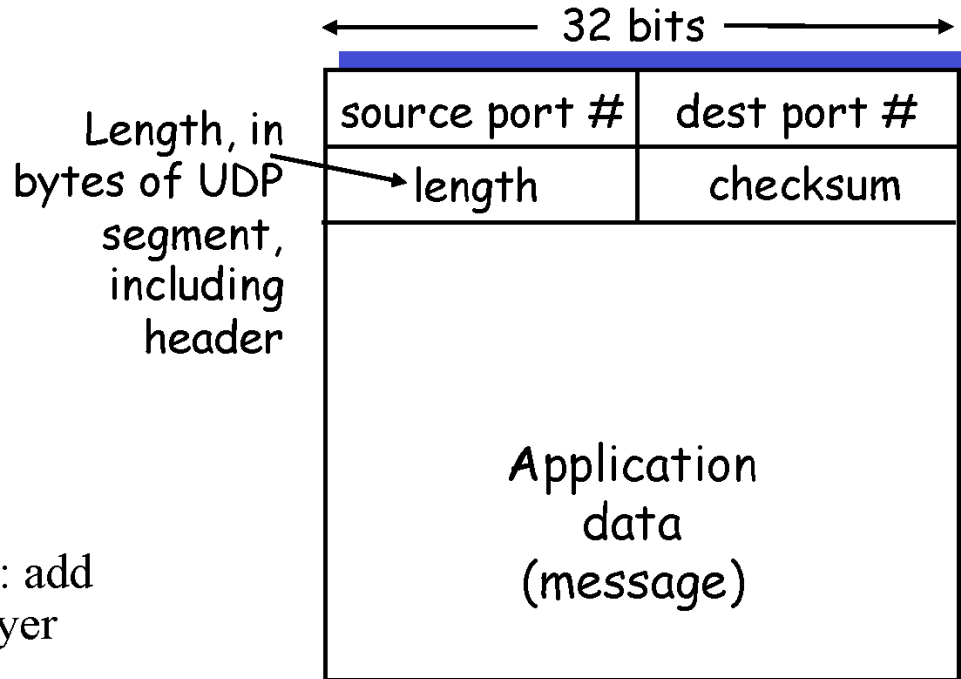
- ❑ “no frills,” “bare bones”
Internet transport protocol
- ❑ “best effort” service, UDP
segments may be:
 - lost
 - delivered out of order to
app
- ❑ *connectionless*:
 - no handshaking between
UDP sender, receiver
 - each UDP segment handled
independently of others

Why is there a UDP?

- ❑ no connection establishment
(which can add delay)
- ❑ simple: no connection state at
sender, receiver
- ❑ small segment header
- ❑ no congestion control: UDP can
blast away as fast as desired

UDP: more

- ❑ often used for streaming multimedia apps
 - loss tolerant
 - rate sensitive
- ❑ other UDP uses
 - DNS
 - SNMP
- ❑ reliable transfer over UDP: add reliability at application layer
 - application-specific error recovery!



UDP segment format