



Chap. 4) Multithreaded Programming

경희대학교 컴퓨터공학과

이 승 룡

■ Heavy-weight

- ✓ A process includes many things:
 - An address space (all the code and data pages)
 - OS resources (e.g., open files) and accounting info.
 - Hardware execution state (PC, SP, registers, etc.)
- ✓ Creating a new process is costly because all of the data structures must be allocated and initialized
 - Linux: over 100 fields in `task_struct` (excluding page tables, etc.)
- ✓ Inter-process communication is costly, since it must usually go through the OS
 - Overhead of system calls and copying data





Thread Concept: Key Idea

- Separate the concept of a process from its execution state
 - ✓ Process: address space, resources, other general process attributes (e.g., privileges)
 - ✓ Execution state: PC, SP, registers, etc.
 - ✓ This execution state is usually called
 - a thread of control,
 - a thread, or
 - a lightweight process (LWP)



Thread Concept: Key Idea (Cont'd)



What is a Thread?

- A *thread* (or *lightweight process*) is a basic unit of CPU utilization; it consists of:
 - ✓ program counter
 - ✓ register set
 - ✓ stack space

- A thread shares with its peer threads its:
 - ✓ code section
 - ✓ data section
 - ✓ operating-system resources
 - ✓ collectively known as a *task* or process

- A traditional or *heavyweight* process is equal to a task with one thread

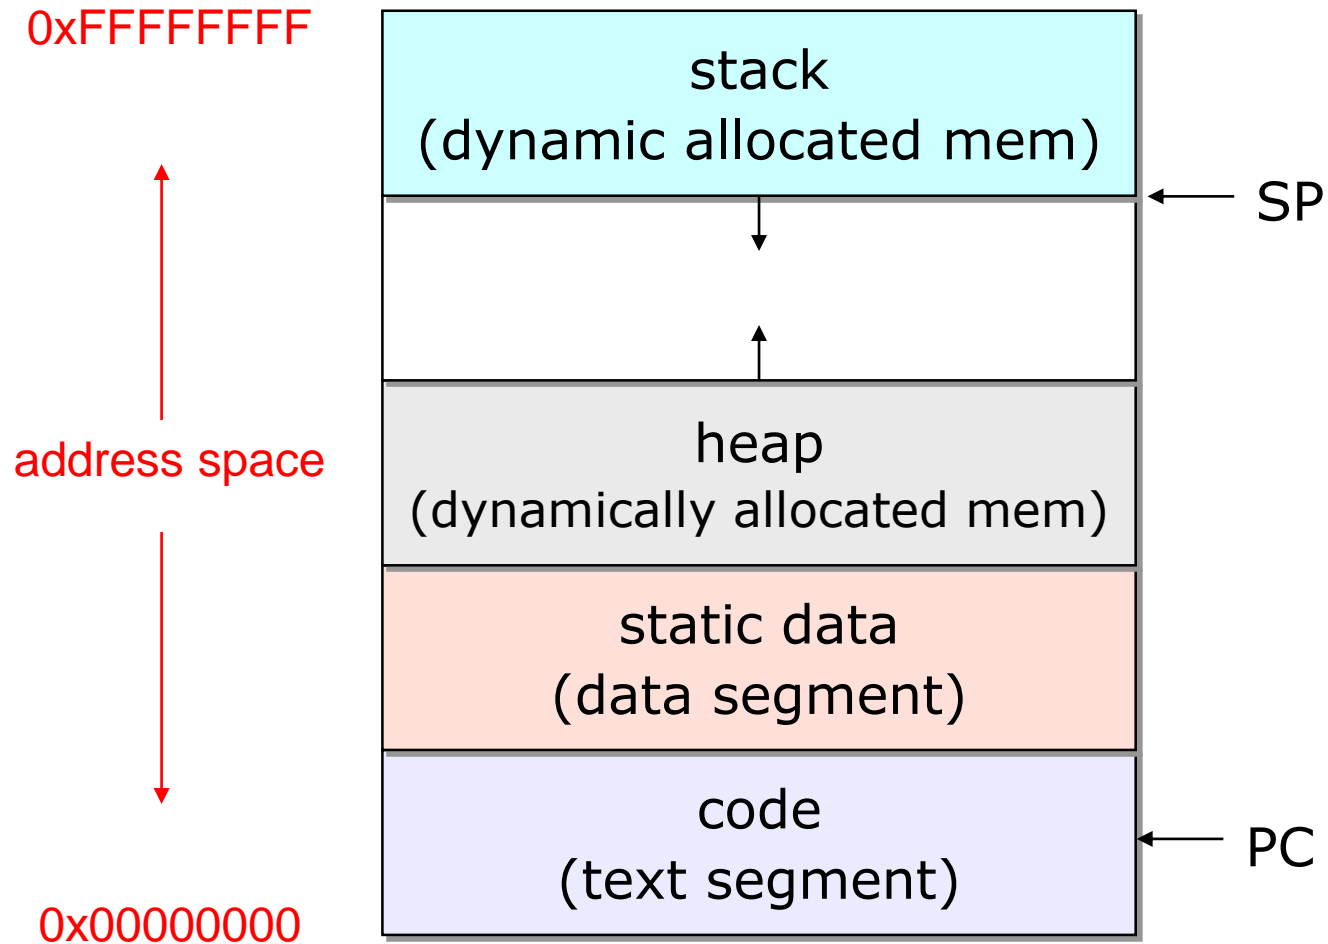


■ Processes vs. Threads

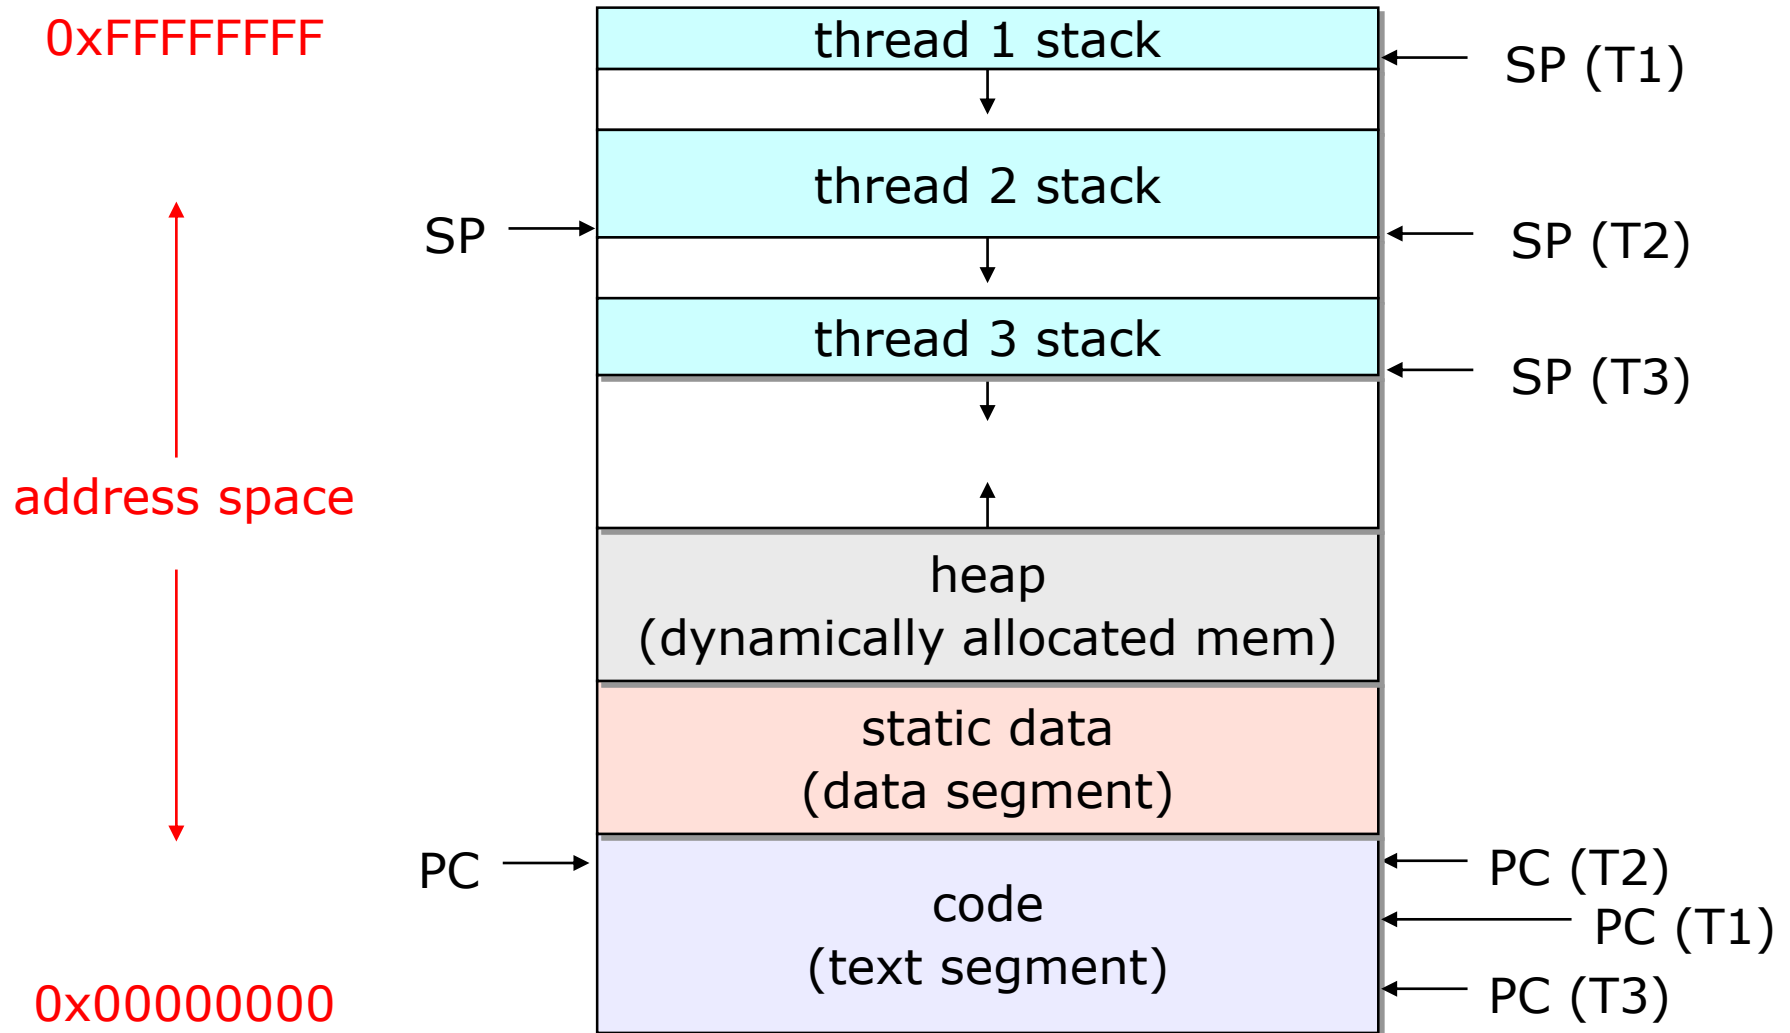
- ✓ A thread is bound to a single process
- ✓ A process, however, can have multiple threads
- ✓ Sharing data between threads is cheap: all see the same address space
- ✓ Threads become the unit of scheduling
- ✓ Processes are now containers in which threads execute
- ✓ Processes become static, threads are the dynamic entities



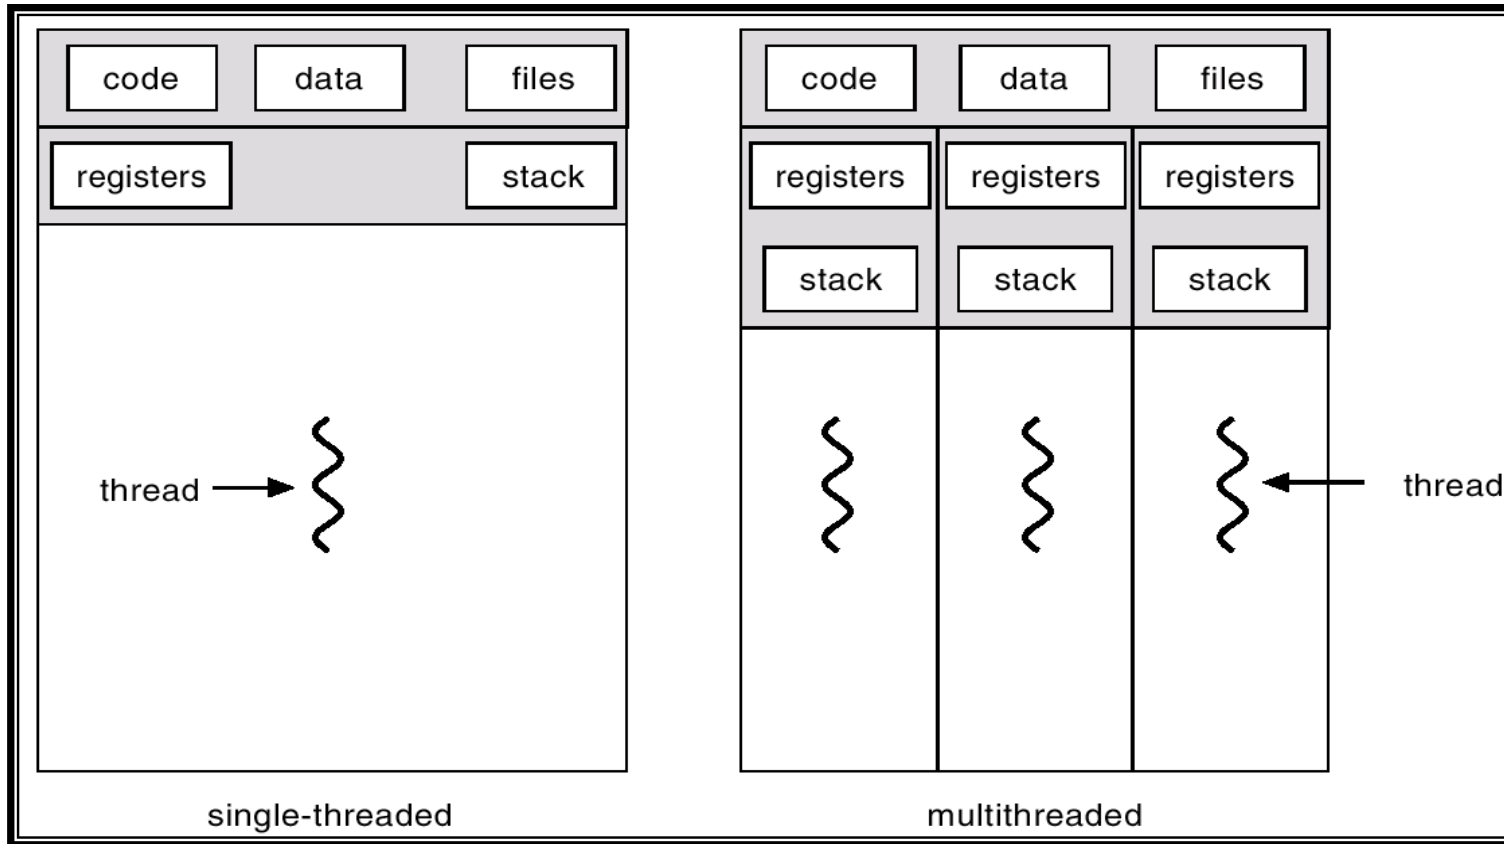
Process Address Space



Address Space with Threads



Single and Multithreaded Processes



Concurrent Servers: Processes

■ Web server example

- ✓ Using fork() to create new processes to handle requests in parallel is overkill for such a simple task.

```
While (1) {  
    int sock = accept();  
    if ((pid = fork()) == 0) {  
        /* Handle client request */  
    } else {  
        /* Close socket */  
    }  
}
```



Concurrent Servers: Threads

■ Using threads

- ✓ We can create a new thread for each request

```
webserver ()
{
    While (1) {
        int sock = accept();
        thread_fork (handle_request, sock);
    }
}
handle_request (int sock)
{
    /* Process request */
    close (sock);
}
```



Benefits

- Responsiveness
- Resource Sharing
- Economy
- Utilization of MP Architectures




User Threads

- Thread management done by user-level threads library

- Examples

- ✓ POSIX *Pthreads*
- ✓ Mach *C-threads*
- ✓ Solaris *threads*




Kernel Threads

■ Supported by the Kernel

- ✓ thread creation and management requires system calls

■ Examples

- ✓ Windows 95/98/NT/2000
- ✓ Solaris
- ✓ Tru64 UNIX
- ✓ BeOS
- ✓ Linux



User-level Threads vs. Kernel-level Threads

■ User-level threads

- ✓ The user-level threads library implements thread operations
- ✓ They are small and fast
- ✓ User-level threads are invisible to the OS
- ✓ OS may make poor decisions
 - e.g. blocking I/O
- ✓ Thread scheduling
 - Non-preemptive scheduling: `yield()`
 - Preemptive scheduling: timer through signal

■ Kernel-level threads

- ✓ All thread operations are implemented in the kernel
- ✓ The OS schedules all of the threads in a system
- ✓ Kernel threads are cheaper than processes
- ✓ They can still be too expensive




Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many




Many-to-One

- Many user-level threads mapped to single kernel thread
- Used on systems that do not support kernel threads




One-to-One

- Each user-level thread maps to kernel thread
- Examples
 - ✓ Windows 95/98/NT/2000/XP
 - ✓ OS/2



Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris 2
- Windows NT/2000/XP with the *ThreadFiber* package



Threading Issues

- Semantics of fork() and exec() system calls
 - ✓ Two versions of fork()
- Thread cancellation
 - ✓ Asynchronous cancellation
 - ✓ Deferred cancellation
- Signal handling
 - ✓ To the thread to which the signal applies
 - ✓ To every thread in the process
 - ✓ To certain threads in the process
 - ✓ Assign a specific thread to receive all signals for the process
- Thread pools
 - ✓ Create a number of threads at process startup
- Thread specific data



Pthreads

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems



■ POSIX-style threads

- ✓ Pthreads
- ✓ DCE threads (early version of Pthreads)
- ✓ Unix International (UI) threads (Solaris threads)
 - Sun Solaris 2, SCO Unixware 2

■ Microsoft-style threads

- ✓ Win32 threads
 - Microsoft Windows 98/NT/2000/XP
- ✓ OS/2 threads
 - IBM OS/2



■ Thread creation/termination

```
int pthread_create (pthread_t *tid,  
                  pthread_attr_t *attr,  
                  void *(start_routine)(void *),  
                  void *arg);
```

```
void pthread_exit  (void *retval);
```

```
int pthread_join   (pthread_t tid,  
                  void **thread_return);
```



■ Mutexes

```
int pthread_mutex_init  
    (pthread_mutex_t *mutex,  
     const pthread_mutexattr_t *mattr);
```

```
int pthread_mutex_destroy  
    (pthread_mutex_t *mutex);
```

```
int pthread_mutex_lock  
    (pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock  
    (pthread_mutex_t *mutex);
```



■ Condition variables

```
int pthread_cond_init  
    (pthread_cond_t *cond,  
     const pthread_condattr_t *cattr);
```

```
int pthread_cond_destroy  
    (pthread_cond_t *cond);
```

```
int pthread_cond_wait  
    (pthread_cond_t *cond,  
     pthread_mutex_t *mutex);
```

```
int pthread_cond_signal  
    (pthread_cond_t *cond);
```

```
int pthread_cond_broadcast  
    (pthread_cond_t *cond);
```



■ Thread creation/termination

```
HANDLE CreateThread (lpThreadAttributes, dwStackSize,  
                    lpStartAddress, lpParameter,  
                    dwCreationFlags, lpThreadId);
```

```
void ExitThread (dwExitCode);
```

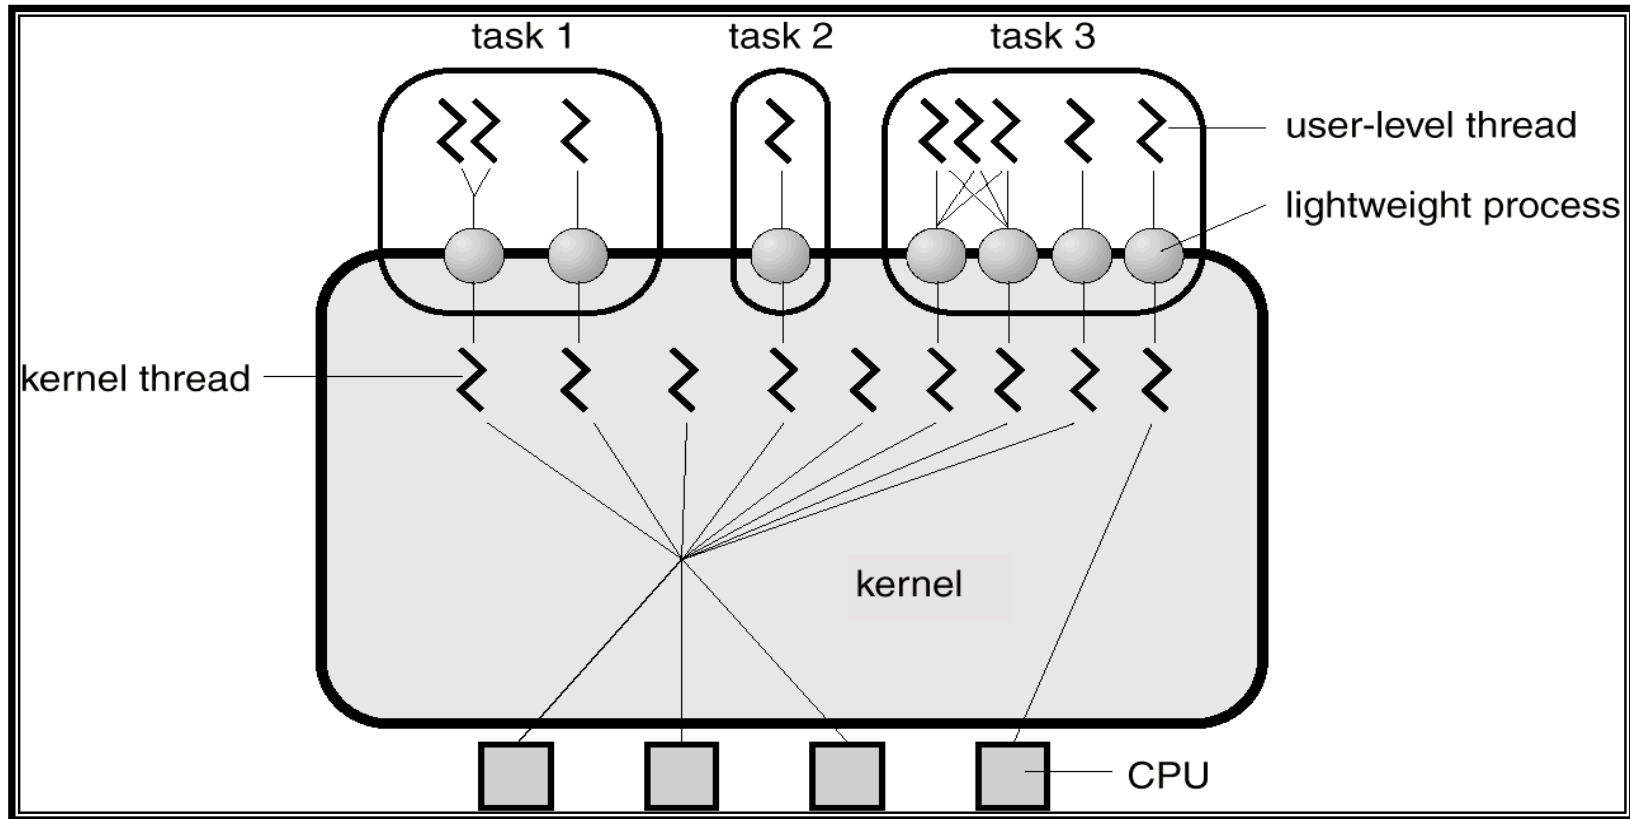


■ Thread creation/termination

Create a new class derived from **Thread** class
Override `run()` method

Create a new class that implements the **Runnable** interface

Solaris 2 Threads



■ LWP (Lightweight Process)

- ✓ A virtual CPU for executing code or system calls
- ✓ Each process contains at least one LWP
- ✓ Each LWP is connected to exactly one kernel-level thread
- ✓ Each LWP is separately dispatched by the kernel, may
 - perform independent system calls
 - incur independent page faults
 - run in parallel on a multiprocessor, etc.
- ✓ The thread library dynamically adjusts the number of LWPs in the pool to ensure the best performance for the application
- ✓ It also “ages” LWPs and deletes them when they are unused for a long time.
- ✓ An LWP is a kernel data structure

■ *For implementing many-to-many model*



Windows XP Threads

- Implements the one-to-one mapping
- Each thread contains
 - ✓ a thread id
 - ✓ register set
 - ✓ separate user and kernel stacks
 - ✓ private data storage area
- Cf) Fibers
 - ✓ Fibers are often called “lightweight” threads
 - ✓ Fibers are invisible to the kernel
 - ✓ Fibers provide a functionality of the many-to-many model



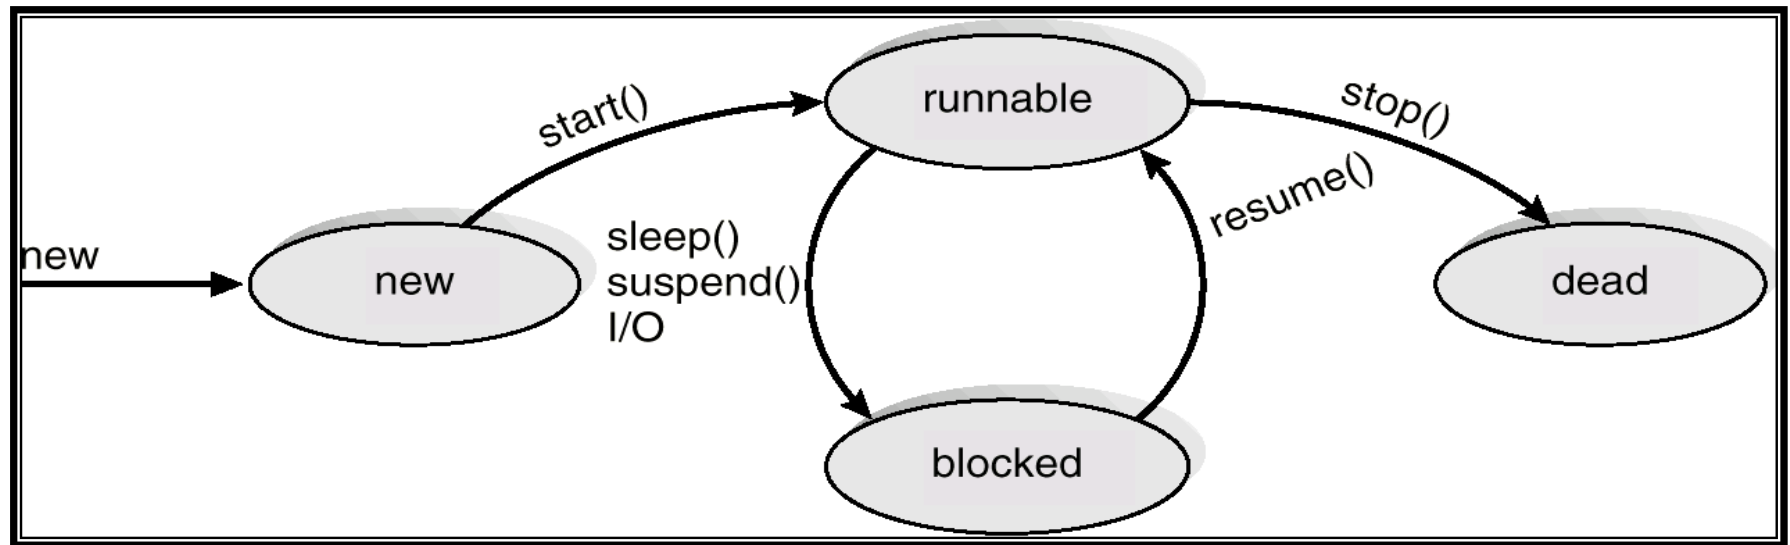
- Linux 2.4 introduces a concept of “thread groups”
 - ✓ *tasks* rather than *threads*
 - ✓ Thread creation is done through clone() system call
 - ✓ Clone() allows a child task to share the address space of the parent task (process)
 - ✓ So, there exist POSIX compatibility problems

- Approaches for POSIX compliance
 - ✓ NPTL (Native POSIX Threading Library) – by RedHat
 - 1:1 model
 - ✓ NGPT (Next Generation POSIX Threading) – by IBM
 - M:N model → Linux 2.6

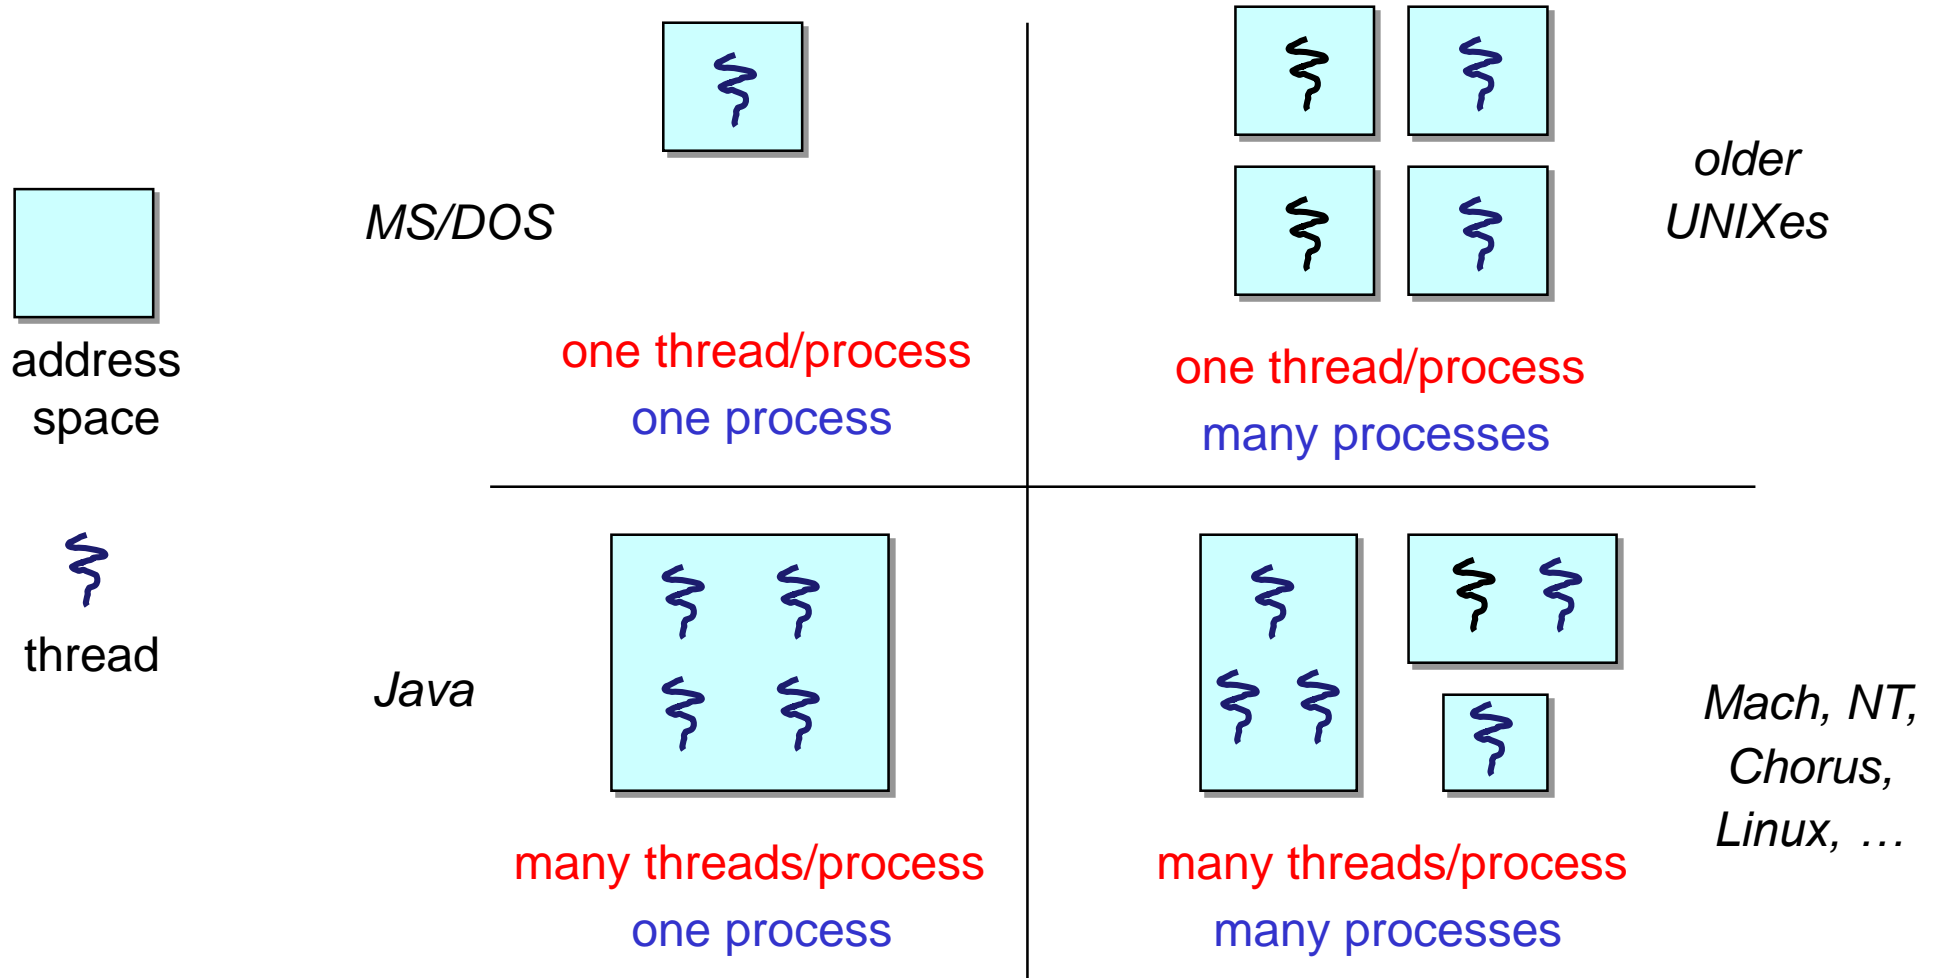


Java Threads

- Java threads may be created by:
 - ✓ Extending Thread class
 - ✓ Implementing the Runnable interface
- Java threads are managed by the JVM
- Java thread states



Threads Design Space



■ Thread concept

- ✓ Separate the concept of a process from its execution state
- ✓ Execution state: PC, SP, registers, etc.

■ Multithreading models

- ✓ User threads to kernel threads mapping
- ✓ Many-to-one
 - User-level threads implementation
- ✓ One-to-one
 - MS-Windows
- ✓ Many-to-many
 - Solaris, Linux

■ Multithreaded programming

- ✓ Unix, Linux: POSIX-style threads (Pthread API)
- ✓ MS-Windows: Win32 threads (Win32 API)
- ✓ Java: Java threads (Thread class)

